

OLS analysis of serially correlated dependent and independent variables

Sandeep Kiran Gudla

May 25, 2019

We know that OLS gives the best linear unbiased estimators when all the [Gauss Markov Assumptions](#) are met. One such assumption is random sampling. Let us see how the estimators behave when this assumption is violated. One of the way to violate this error is creating serial dependencies on dependent and independent variables.

Let us begin by generating dependent and independent variables using the following data generating processes.

$$x_t = x_{t-1,i} + u_{t,i}, u_{t,i} \sim iid\mathcal{N}(0, \sigma^2)$$

for $t = 1, 2, \dots, T$ and $i = 1, 2, \dots, n$.

Let's use a similar process to generate another feature, say $\{y_t\}$

$$y_t = y_{t-1} + e_t, e_t \sim iid\mathcal{N}(0, \sigma^2)$$

for $t = 1, 2, \dots, T$

```
#importing the necessary packages
import numpy as np
import pandas as pd
import random
from statsmodels import api as sm
import matplotlib.pyplot as plt
import scipy as sc
```

Let's generate $x_{t,i}$ for $t = 1$ and $i = 1$, $x_1 = x_0 + u_{1,i}, u_{1,i} \sim iid\mathcal{N}(0, \sigma^2)$ and y_t for $t = 1$, $y_1 = y_0 + e_1, e_1 \sim iid\mathcal{N}(0, \sigma^2)$


```
t=1
e=[]
y=[]
u=[]
x=[]
#initial y gen
y.append(random.normalvariate(4,1))
#initial x gen
x.append(random.normalvariate(10,1))
#generate errors
simn = 1000
```

```

random.seed(2)
for isim in range(0,simn):
    #error e gen
    e.append(random.normalvariate(0,1))
    #error u gen
    u.append(random.normalvariate(0,1))
for isim in range(1,simn):
    #generate y
    y.append(y[isim-t]+e[isim])
    #generate x
    x.append(x[isim-t]+u[isim])
x=np.array(x).reshape(-1,1)
y=np.array(y)
ones = np.ones((len(x),1))
x = np.hstack((ones,x))

# the histogram of the error
n, bins, patches = plt.hist(e, 10, density=True, facecolor='g', alpha=0.75)
plt.grid(True)
plt.show()

```




Here, we can see that the error is normally distributed about zero.

Let's analyze y

```

# the histogram of the y
n, bins, patches = plt.hist(y, 5, density=True, facecolor='g', alpha=0.75)
plt.grid(True)
plt.show()


```




We can see that y is not normally distributed.

Similarly, let's see the histograms of u and x.

```
# the histogram of the u  
n, bins, patches = plt.hist(u, 10, density=True, facecolor='g', alpha=0.75)  
plt.grid(True)  
plt.show()
```



```
# the histogram of the x
n, bins, patches = plt.hist(x[:,1], 5, density=True, facecolor='g', alpha=0.75)
plt.grid(True)
plt.show()
```



Let's build the following model using OLS

$$y_t = \beta_0 + \beta_1 x_t + \epsilon_t$$

```
# regression of x on y
model1 = sm.OLS(y,x).fit()
print(model1.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          y      R-squared:          0.732
Model:                  OLS    Adj. R-squared:       0.732
Method:                 Least Squares  F-statistic:       2730.
Date:                   Tue, 21 May 2019  Prob (F-statistic):  7.47e-288
Time:                   17:03:52  Log-Likelihood:    -2860.2
No. Observations:      1000      AIC:              5724.
Df Residuals:          998       BIC:              5734.
Df Model:               1
Covariance Type:       nonrobust
=====
```


	coef	std err	t	P> t	[0.025	0.975]
-----	-----	-----	-----	-----	-----	-----
const	8.7965	0.167	52.775	0.000	8.469	9.124
x1	-0.3762	0.007	-52.249	0.000	-0.390	-0.362
=====	=====	=====	=====	=====	=====	=====
Omnibus:		6.544	Durbin-Watson:			0.061
Prob(Omnibus):		0.038	Jarque-Bera (JB):			5.690
Skew:		-0.118	Prob(JB):			0.0581
Kurtosis:		2.715	Cond. No.			28.9
=====	=====	=====	=====	=====	=====	=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

We can see that the [Durbin-Watson](#) score is 0.061 which suggests that there is positive auto correlation

```
# the histogram of the data
residuals=np.array(model1.resid).reshape(-1,1)
n, bins, patches = plt.hist(residuals, 20, density=True, facecolor='g', alpha=0.75)
#plt.axis([-1, 1, 0, 20])
plt.grid(True)
plt.show()
```



Here, we can see that the error terms are not normally distributed, there for the assumptions are violated and the OLS estimators are biased.

Lets look into the correlation of error terms.

```
model1 = sm.OLS(y,x).fit(cov_type='HCO')
print(model1.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  y      R-squared:                0.732
Model:                        OLS      Adj. R-squared:           0.732
Method:                    Least Squares  F-statistic:                3141.
Date:                Tue, 21 May 2019  Prob (F-statistic):        1.58e-310
Time:                17:03:56      Log-Likelihood:            -2860.2
No. Observations:          1000      AIC:                      5724.
Df Residuals:              998      BIC:                      5734.
Df Model:                  1
Covariance Type:            HCO
=====

               coef      std err          z      P>|z|      [0.025      0.975]
-----
const          8.7965      0.153      57.330      0.000      8.496      9.097
x1          -0.3762      0.007     -56.045      0.000     -0.389     -0.363
=====

Omnibus:                 6.544  Durbin-Watson:           0.061
Prob(Omnibus):            0.038  Jarque-Bera (JB):        5.690
Skew:                    -0.118  Prob(JB):               0.0581
Kurtosis:                 2.715  Cond. No.               28.9
=====


```

Warnings:

[1] Standard Errors are heteroscedasticity robust (HCO)

Let's plot residual plots to get a better understanding of the error distribution.

```
plt.scatter(x[:,1],residuals.flatten())
```



We can see that there is heterogeneity in error terms. This leads to bias in our estimators and therefore, the OLS estimators are no longer BLUE (best linear unbiased estimators).

```
sc.stats.levene(y,x[:,1])
```

```
LeveneResult - statistic = 663.1214037868205, pvalue=1.6141243531100978e-126
```

```
sc.stats.ttest_ind(y,x[:,1])
```

```
Ttest_indResult - statistic = 43.294390393344614, pvalue=2.0469050973291668e-289
```

Based on the [Levene Test](#) and [Welch's t-test](#), the null hypothesis that the random sampling distribution has equal variance can be rejected.

Let's analyze the case for Multivariate Linear Regressions

```
t=1
e=[]
y=[]
u=[]
x=[]
#initial y gen
n = 2
y.append(random.normalvariate(4,1))
#initial x gen
x.append(np.random.normal(10,1,n))
#generate errors
```

```

simn = 1000
random.seed(2)
for isim in range(0,simn):
    #error e gen
    e.append(random.normalvariate(0,1))
    #error u gen
    u.append(np.random.normal(0,1,n))
for isim in range(1,simn):
    #generate y
    y.append(y[isim-t]+e[isim])
    #generate x
    x.append(x[isim-t]+u[isim])
x=np.array(x).reshape(-1,n)
y=np.array(y)
ones = np.ones((len(x),1))
x = np.hstack((ones,x))

# regression of x on y
modell = sm.OLS(y,x).fit()
print(modell.summary())

```

OLS Regression Results

Dep. Variable:	y	R-squared:	0.612			
Model:	OLS	Adj. R-squared:	0.612			
Method:	Least Squares	F-statistic:	787.7			
Date:	Tue, 21 May 2019	Prob (F-statistic):	6.15e-206			
Time:	20:45:46	Log-Likelihood:	-3221.0			
No. Observations:	1000	AIC:	6448.			
Df Residuals:	997	BIC:	6463.			
Df Model:	2					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	12.0352	0.426	28.281	0.000	11.200	12.870
x1	-0.2656	0.031	-8.548	0.000	-0.327	-0.205
x2	-0.7065	0.022	-31.722	0.000	-0.750	-0.663
=====						
Omnibus:	14.657	Durbin-Watson:	0.040			
Prob(Omnibus):	0.001	Jarque-Bera (JB):	13.561			
Skew:	0.239	Prob(JB):	0.00114			
Kurtosis:	2.689	Cond. No.	41.4			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


```

model1 = sm.OLS(y,x).fit(cov_type='HCO')
print(model1.summary())

```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          y      R-squared:          0.612
Model:                  OLS    Adj. R-squared:      0.612
Method:                 Least Squares  F-statistic:      828.4
Date:                   Tue, 21 May 2019  Prob (F-statistic):  1.11e-212
Time:                   20:48:45   Log-Likelihood:    -3221.0
No. Observations:      1000      AIC:              6448.
Df Residuals:          997      BIC:              6463.
Df Model:               2
Covariance Type:       HCO
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	12.0352	0.329	36.594	0.000	11.391	12.680
x1	-0.2656	0.033	-7.939	0.000	-0.331	-0.200
x2	-0.7065	0.019	-36.516	0.000	-0.744	-0.669

```

=====
Omnibus:                 14.657   Durbin-Watson:          0.040
Prob(Omnibus):            0.001   Jarque-Bera (JB):       13.561
Skew:                     0.239   Prob(JB):               0.00114
Kurtosis:                 2.689   Cond. No.                41.4
=====

```

```


Warnings:
[1] Standard Errors are heteroscedasticity robust (HCO)

```


```

plt.scatter(x[:,1],residuals.flatten())

```



```
plt.scatter(x[:,2],residuals.flatten())
```



```
sc.stats.levene(y,x[:,1],x[:,2])
```

```
LeveneResult - statistic = 160.18648446129512, pvalue=8.031994802470273e-67
```


```
[sc.stats.ttest_ind(y,x[:,1]),sc.stats.ttest_ind(y,x[:,2])]
```

```
Ttest_indResult - statistic = 6.957613621196703, pvalue=4.673644578948617e-12,
```

```
Ttest_indResult - statistic = -34.82903392988044, pvalue=4.1258493384478805e-208
```

Now let's analyze how the number of samples effect the overall R-square

```
t=1
e=[]
y=[]
u=[]
x=[]
#initial y gen
y.append(random.normalvariate(4,1))
#initial x gen
x.append(random.normalvariate(10,1))
#generate errors
simn = 1000
random.seed(2)
R_sq = [0]
for isim in range(0,simn):
    #error e gen
    e.append(random.normalvariate(0,1))
    #error u gen
    u.append(random.normalvariate(0,1))
for isim in range(1,simn):
    #generate y
    y.append(y[isim-t]+e[isim])
    #generate x
    x.append(x[isim-t]+u[isim])
    x_ar = np.array(x).reshape(-1,1)
    ones = np.ones((len(x),1))
    x_ar = np.hstack((ones,x_ar))
    y_ar=np.array(y)
    model1 = sm.OLS(y_ar,x_ar).fit()
    R_sq.append(model1.rsquared)
plt.plot(range(1,simn),R_sq[1:])
```



```


t=1
e=[]
y=[]
u=[]
x=[]
#initial y gen
n = 2
y.append(random.normalvariate(4,1))
#initial x gen
x.append(np.random.normal(10,1,n))
#generate errors
simn = 1000
random.seed(2)
R_sq = [0]
for isim in range(0,simn):
    #error e gen
    e.append(random.normalvariate(0,1))
    #error u gen
    u.append(np.random.normal(0,1,n))
for isim in range(1,simn):
    #generate y
    y.append(y[isim-t]+e[isim])
    #generate x
    x.append(x[isim-t]+u[isim])
    x_ar = np.array(x).reshape(-1,n)
    ones = np.ones((len(x),1))

```

```

x_ar = np.hstack((ones,x_ar))
y_ar=np.array(y)
model1 = sm.OLS(y_ar,x_ar).fit()
R_sq.append(model1.rsquared)
plt.plot(range(1,simn),R_sq[1:])

```



Lets start analyzing for different time lags starting from $t = 1, 2, 3, \dots, 10$

```

e=[]
y=[]
u=[]
x=[]
#initial y gen
n = 1
y.append(random.normalvariate(4,1))
#initial x gen
x.append(np.random.normal(10,1,n))
#generate errors
simn = 1000
outp = []
time = []
R_sq = []
random.seed(2)
T = [1,5,10,15,20,25,30,35,40,45,50]
for t in T:
    for isim in range(0,simn):
        #error e gen


```

```


e.append(random.normalvariate(0,1))
#error u gen
u.append(np.random.normal(0,1,n))
for isim in range(1,simn):
    #generate y
    y.append(y[isim-t]+e[isim])
    #generate x
    x.append(x[isim-t]+u[isim])
    x_ar = np.array(x).reshape(-1,n)
    ones = np.ones((len(x),1))
    x_ar = np.hstack((ones,x_ar))
    y_ar=np.array(y)
    model1 = sm.OLS(y_ar,x_ar).fit()
    outp.append(model1.rsquared)
    #time.append(t)
print("time lags - ",t)
R_sq.append(model1.rsquared)
plt.plot(range(1,simn),outp)
plt.show()
outp=[]

```


time lags - 1




time lags - 5




time lags - 10




time lags - 15




time lags - 20




time lags - 25




time lags - 30




time lags - 35




time lags - 40



time lags - 45



time lags - 50



Now consider the case where $n = 1, \dots, 10$ and $t \in [1, 5, 10, 15, 20, 25, 30, 35, 40, 45, 50]$


```

simn = 1000
outp = []
time = []
R_sq = []
random.seed(2)
T = [1,5,10,15,20,25,30,35,40,45,50]
for n in range(1,10):
    #generate errors
    e=[]
    y=[]
    u=[]
    x=[]
    #initial y gen
    y.append(random.normalvariate(4,1))
    #initial x gen
    x.append(np.random.normal(10,1,n))
    for t in T:
        for isim in range(0,simn):
            #error e gen
            e.append(random.normalvariate(0,1))
            #error u gen
            u.append(np.random.normal(0,1,n))
        for isim in range(1,simn):
            #generate y
            y.append(y[isim-t]+e[isim])
            #generate x
            x.append(x[isim-t]+u[isim])
            x_ar = np.array(x).reshape(-1,n)
            ones = np.ones((len(x),1))
            x_ar = np.hstack((ones,x_ar))
            y_ar=np.array(y)
            model1 = sm.OLS(y_ar,x_ar).fit()
            outp.append(model1.rsquared)
            #time.append(t)
        print("number of parameters - ",n)
        print("time lags - ",t)
        R_sq.append(model1.rsquared)
        plt.plot(range(1,simn),outp)
        plt.show()
        outp=[]


```

number of parameters - 1


time lags - 1




number of parameters - 1
time lags - 5




number of parameters - 1
time lags - 10




number of parameters - 1
time lags - 15




number of parameters - 1
time lags - 20




number of parameters - 1
time lags - 25




number of parameters - 1
time lags - 30




number of parameters - 1
time lags - 35




number of parameters - 1
time lags - 40




number of parameters - 1
time lags - 45




number of parameters - 1
time lags - 50




number of parameters - 2
time lags - 1




number of parameters - 2
time lags - 5




number of parameters - 2
time lags - 10




number of parameters - 2
time lags - 15




number of parameters - 2
time lags - 20




number of parameters - 2
time lags - 25




number of parameters - 2
time lags - 30




number of parameters - 2
time lags - 35




number of parameters - 2
time lags - 40




number of parameters - 2
time lags - 45




number of parameters - 2
time lags - 50




number of parameters - 3
time lags - 1




number of parameters - 3
time lags - 5




number of parameters - 3
time lags - 10




number of parameters - 3
time lags - 15




number of parameters - 3
time lags - 20




number of parameters - 3
time lags - 25




number of parameters - 3
time lags - 30




number of parameters - 3
time lags - 35




number of parameters - 3
time lags - 40




number of parameters - 3
time lags - 45




number of parameters - 3
time lags - 50




number of parameters - 4
time lags - 1




number of parameters - 4
time lags - 5




number of parameters - 4
time lags - 10




number of parameters - 4
time lags - 15




number of parameters - 4
time lags - 20




number of parameters - 4
time lags - 25




number of parameters - 4
time lags - 30




number of parameters - 4
time lags - 35




number of parameters - 4
time lags - 40



number of parameters - 4
time lags - 45



number of parameters - 4
time lags - 50



Similar plots are observed for higher order of t and N ; where t = time - lags, N = number of parameters.

Conclusion We saw that, the OLS output of serially correlated data gives spurious results, the OLS estimators are no longer best estimators because the random sampling assumption is violated.

In coming posts, we will discuss on how to fix the serial correlation issue in OLS.